Databender

FREE GUIDE

The Economics Changed

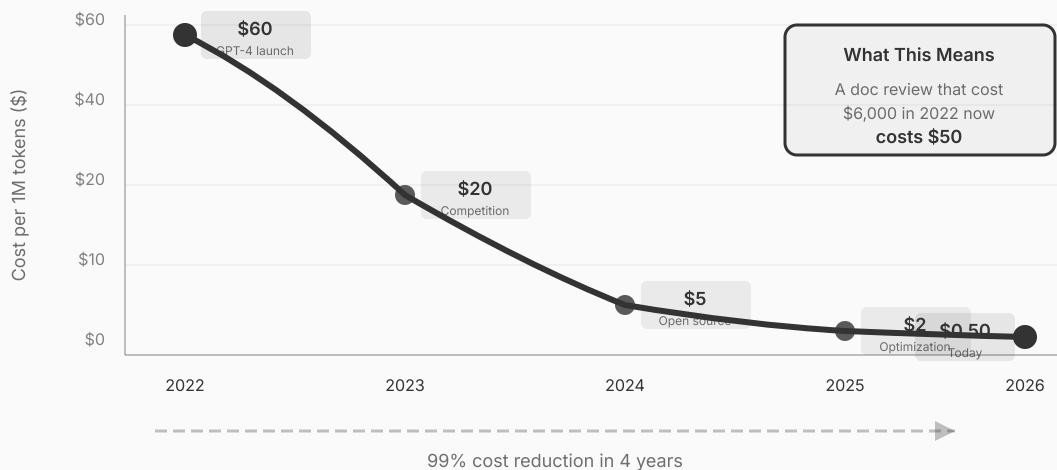What Used to Cost $200K Now Costs $30-50K

databender.co | Data & AI Consulting for Growing Businesses

80%

Cost Reduction

Custom legal technology projects cost 80% less than they did just two years ago.

The Economics of AI: Rapidly Declining Costs

Cost per 1M tokens ($)

- $60 — GPT-4 launch
- $20 — Competition
- $5 — Open source
- $2 — Optimization
- $0.50 — Today

What This Means
A doc review that cost $6,000 in 2022 now **costs $50**

2022 · 2023 · 2024 · 2025 · 2026

99% cost reduction in 4 years

Two years ago, a managing partner asked us about building a document search system for her firm. We quoted $180K. She laughed, said "maybe someday," and went back to watching associates bill three hours to find precedent the firm had written ten years earlier.

We talked to her again last month. Same project. $42K. Six weeks.

She stopped laughing. She started building.

Why Prices Dropped 80%

Here's what changed: we stopped building from scratch.

In 2022, a document search system meant writing custom code for text extraction, building search algorithms, designing interfaces, testing everything. Four developers for six months. $300K minimum, and that's before anyone actually used it.

Now? The hard parts come pre-built. Language models that understand legal concepts. Search systems that find meaning, not just keywords. Interface patterns that work. A senior consultant assembles these pieces instead of inventing them.

The analogy: building a house from trees you cut down versus building from lumber and prefab components. Same house. Fraction of the time. The skill shifted from carpentry to architecture.

What This Means for Your Firm

Projects that weren't worth considering are now worth considering.

Making your document repository searchable in plain English? That used to be a $150K project. Now it's $35-50K. Your associates ask "what's our standard approach to California non-competes?" and get a useful answer in seconds.

Building a system that captures knowledge from retiring partners? Enterprise vendors quote six figures. We've built them for $40K.

Contract review automation tailored to how your firm actually works? Not the one-size-fits-all SaaS tool. Custom logic for your practice areas. $30-50K, done in weeks.

Real Project Costs

The document intelligence build: $45,000

A real estate practice in Phoenix. Fifty attorneys, twenty years of transaction documents scattered across network drives and an aging DMS. Associates spent half their research time just finding things.

We indexed everything. Contracts, memos, closing documents, correspondence. Six weeks later, an associate types "what's our standard approach to Arizona water rights easements" and gets three relevant precedents with citations. The partner who handled those deals retired in 2019. His knowledge didn't leave with him.

> **Knowledge preservation: $38,000.** Two founding partners retiring within eighteen months. The firm captured three decades of client intelligence, deal patterns, and relationship context before it walked out the door.

> **Contract review: $52,000.** Corporate group reviewing 200+ vendor agreements monthly. Custom extraction for their clause library. Review time dropped 60%.

> **Research acceleration: $28,000.** Litigation boutique integrated AI directly into their document management system. Research that took hours now takes minutes.

What We Learned From Other Industries

Here's something we didn't expect: the patterns are nearly identical across industries.

We've built document intelligence systems for healthcare organizations and manufacturers. The problems look different on the surface. Protocols vs. contracts. Specs vs. briefs. But the underlying challenge is the same: decades of institutional knowledge trapped in documents nobody can find.

What worked for a medical device distributor worked for a law firm. What worked for a manufacturer's quality documentation worked for litigation research. The solutions transfer. The lessons compound.

This matters because you're not our first build. You're benefiting from patterns we've refined across dozens of implementations.

How Builds Actually Work

The first two weeks are about understanding. What documents exist? What questions do associates actually ask? Where does research time disappear? We watch people work. We find the friction.

Weeks three and four are construction. This is where pre-built components matter. We're not writing a search engine from scratch. We're configuring one that already works, training it on your firm's specific language and patterns.

The final stretch is refinement. Real associates using the system. Finding edge cases. Tuning results. By week six, attorneys are using it for actual matters.

We're not going to pretend every build fits this timeline perfectly. Complex integrations take longer. But six to eight weeks is realistic for most firms, and that's a fraction of what enterprise deployments require.

Why You Own the Result

SaaS tools rent you capability. Per seat, per month, forever. The vendor holds your data and your future.

Custom builds work differently. You own the code. It runs on your servers. Add users without adding costs. Switch consultants without losing your system.

Five-year comparison: SaaS at $80/user/month for 50 users costs $240,000. Custom build at $50,000 plus updates costs $75,000. That's $165,000 in savings, and you own something at the end.

Ethics Is Easier Than You Think

Cloud AI creates ethics headaches. Client data in someone else's servers. Privacy policies that change. Security you can't verify.

Custom AI on your servers sidesteps this entirely. Does client data leave your building? No. Can you audit everything? Yes. Can you explain exactly where data goes to your ethics committee? Completely.

databender.co

Ethics committees that spend months evaluating cloud vendors often approve on-premise deployments in weeks. There's simply less to evaluate when nothing leaves your control.

Where Firms Go Wrong

Most firms that "explore AI" are still exploring two years later. The pattern is predictable.

IT committee forms. Vendors present. Security review begins. Six months pass. The committee recommends a pilot. The pilot takes three months to scope. By the time anything launches, eighteen months have evaporated and the associates who were supposed to use it have moved to other firms.

The firms actually deploying AI skip this theater. They start with one specific problem. They choose on-premise to sidestep the cloud security debate entirely. They build something small, prove it works, then expand.

The difference isn't budget or sophistication. It's recognizing that the biggest risk isn't moving too fast. It's waiting while competitors build advantages you'll spend years trying to match.

The First-Mover Advantage

In twelve months, managing partners who've been "monitoring AI developments" will finally decide to act. They'll start where you could start today.

But you won't be where you are today. You'll have a system that's been learning your firm's patterns for a year. Associates who've internalized new workflows. A knowledge base that grows with every matter.

The firms that moved first on legal research tools in the 2010s didn't just save money. They trained a generation of associates to work differently. The ones who waited spent years catching up.

This is that moment again.

Where to Start

One managing partner put it this way: "I'd spent two years waiting for the 'right time' to look at AI. When I finally saw the numbers, I realized I'd already spent more than a custom build would have cost, just paying associates to do work a system could have handled."

Pick one problem. Document search. Knowledge capture. Contract review. Get a real quote for your specific situation.

If the math works, build it. If it doesn't, you've lost nothing but an hour of conversation.

The economics changed. The firms that noticed are already ahead.

Ready to see what's possible at today's prices? Schedule a conversation or explore our legal-specific solutions.